

# mongoDB

Bộ môn: Kỹ Thuật Phần Mềm

Giáo viên: Trần Thế Trung.  
Email: [tranthewtrung@iu.edu.vn](mailto:tranthewtrung@iu.edu.vn)



# MongoDB Basic Cluster Administration



1. Giới thiệu về **Mongod**.
2. Giới thiệu về **Replication**.
3. Giới thiệu về **Sharding**.

# 1. mongod

# Learning Objectives

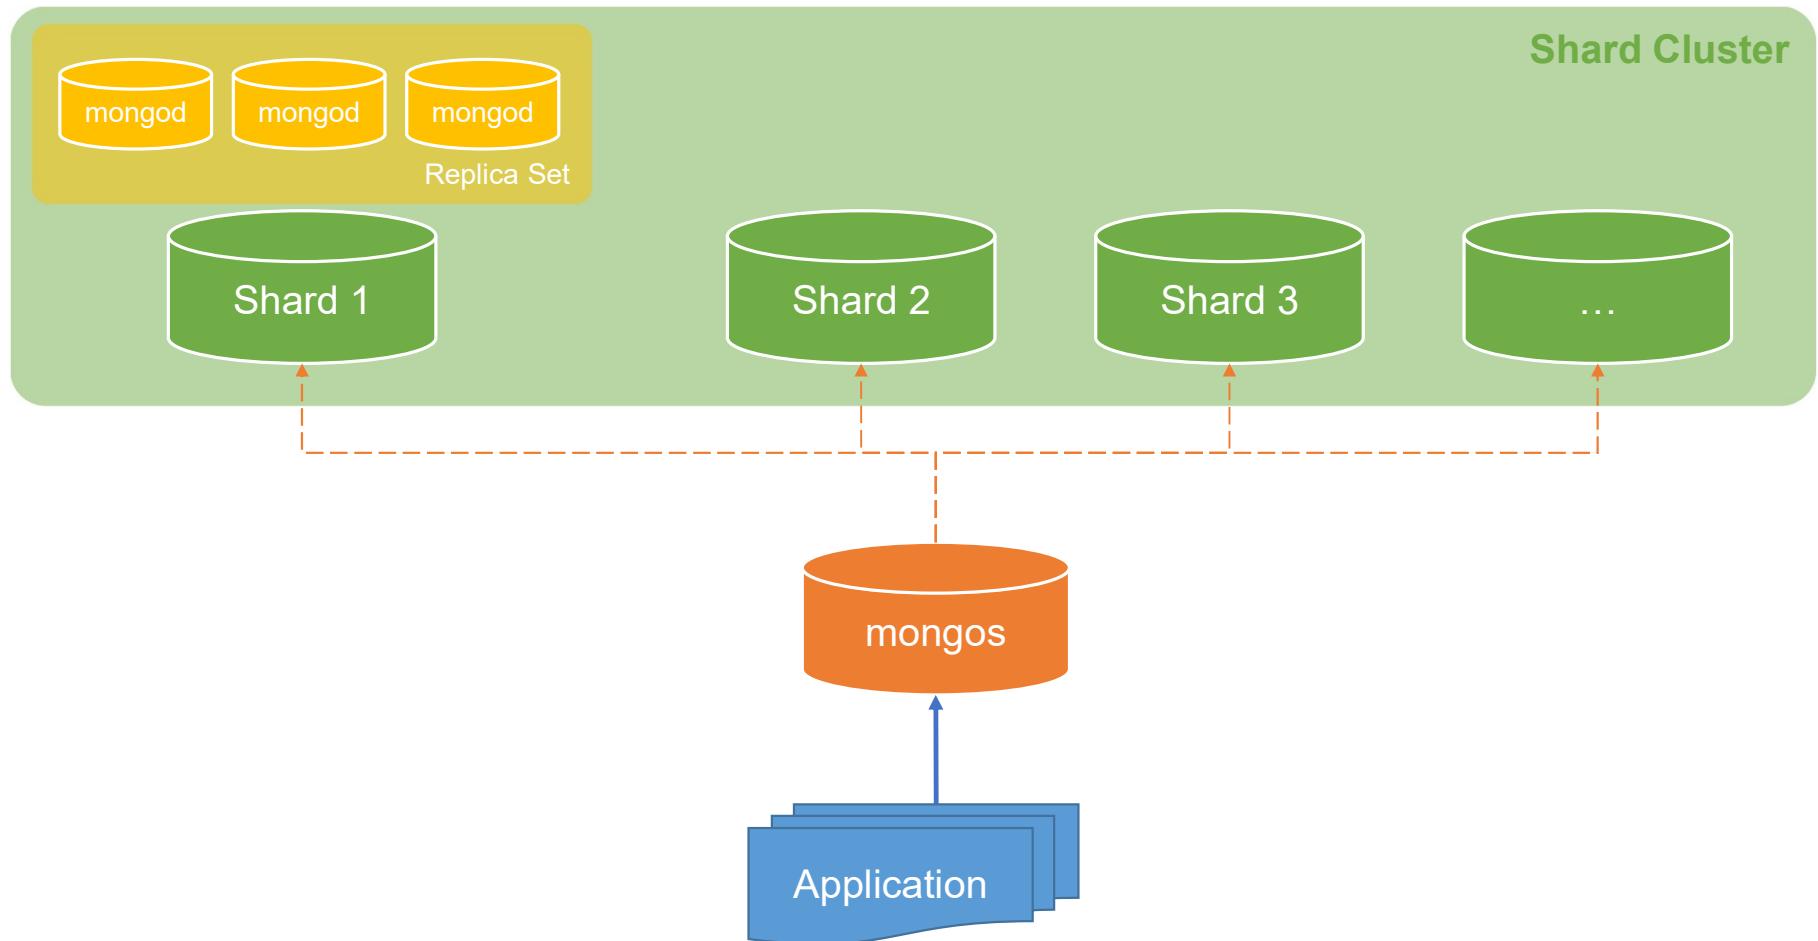
- What `mongod` is?
- How to communicate with `mongod` ?
- Default configuration for `mongod`.

# What mongod is

- mongod is the main **daemon process** for MongoDB.
- The **core server** of the database, handling connections, requests, and most importantly, persisting your data.
- MongoDB deployment may consist of **more than one server**. Our data may be distributed in a **replica set** or across a **sharded cluster**.
- We run a **separate mongod process** for each server.



# What mongod is

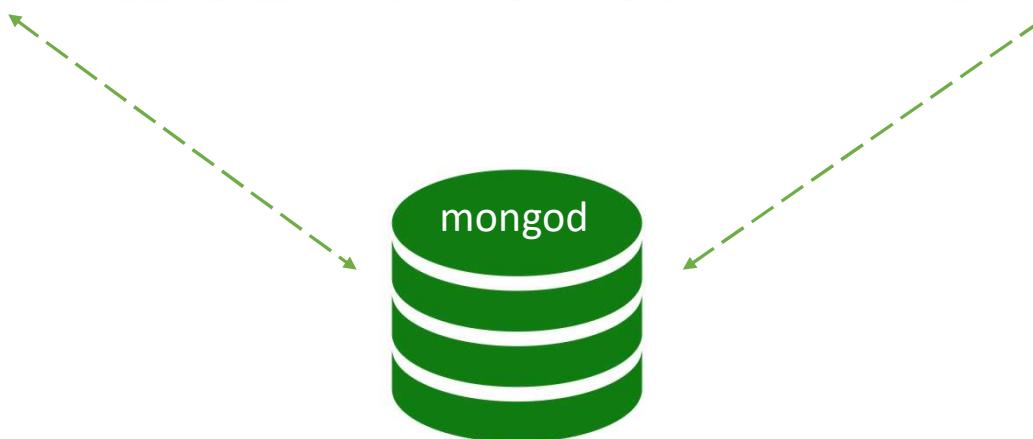


# What mongod is

- When we launch mongod, we're essentially starting up a new database. But we don't interact with the mongod process directly. Instead, we use a database client to communicate with mongod (*mongosh*, *mongo*).

```
C:\WINDOWS\system32\cmd.exe - mongo
To enable free monitoring, run the following command: db.en
ableFreeMonitoring()
To permanently disable this reminder, run the following com
mand: db.disableFreeMonitoring()
-->
```

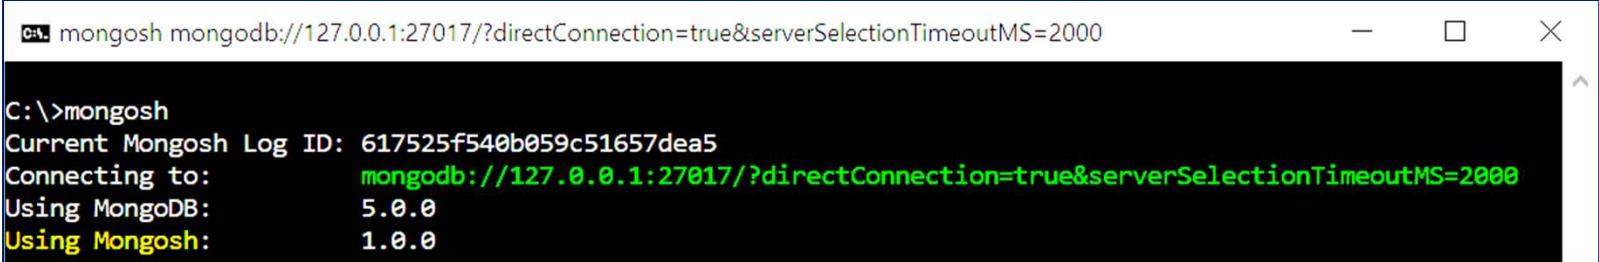
```
C:\Users\ZENBOOK>mongosh
Current Mongosh Log ID: 6220444e20ed9f539dd0e6fe
Connecting to:      mongodb://127.0.0.1:27017/?directConnection
=true&serverSelectionTimeoutMS=2000
Using MongoDB:     5.0.5
Using Mongosh:     1.1.7
```



# What mongod is

Default Configuration:

- The port mongod listens on will default to **27017**.
- The default dbpath is **/data/db** (*this folder should be available when we run mongod*).
- Bind to localhost by default (127.0.0.1).
- Authentication is turned off by default, so clients are not required to authenticate before accessing the database.



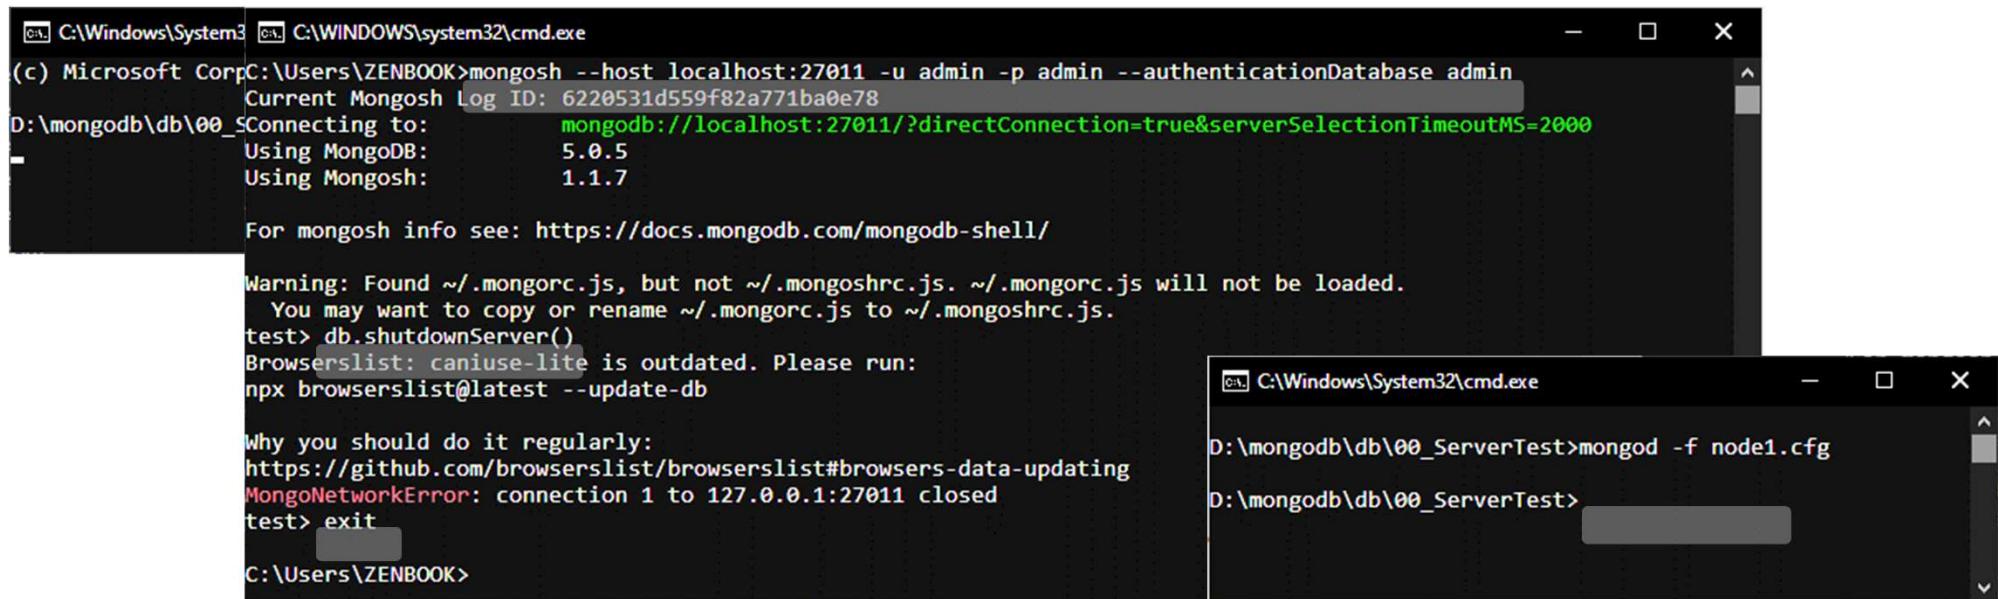
```
mongosh mongodb://127.0.0.1:27017/?directConnection=true&serverSelectionTimeoutMS=2000
C:\>mongosh
Current Mongosh Log ID: 617525f540b059c51657dea5
Connecting to:      mongodb://127.0.0.1:27017/?directConnection=true&serverSelectionTimeoutMS=2000
Using MongoDB:    5.0.0
Using Mongosh:   1.0.0
```

# What mongod is

To start up a mongod process': mongod ↴

To shutdown mongod from mongo shell (mongosh):

- use admin ↴
- db.shutdownServer() ↴
- exit ↴ (exit mongosh)



The screenshot shows a Windows Command Prompt window titled 'C:\Windows\System32 C:\WINDOWS\system32\cmd.exe'. It displays the following terminal session:

```
(c) Microsoft Corp C:\Users\ZENBOOK>mongosh --host localhost:27011 -u admin -p admin --authenticationDatabase admin
Current Mongosh Log ID: 6220531d559f82a771ba0e78
D:\mongodb\db\00_ServerTest> Connecting to: mongodb://localhost:27011/?directConnection=true&serverSelectionTimeoutMS=2000
Using MongoDB: 5.0.5
Using Mongosh: 1.1.7

For mongosh info see: https://docs.mongodb.com/mongodb-shell/

Warning: Found ~/.mongorc.js, but not ~/.mongoshrc.js. ~/.mongorc.js will not be loaded.
You may want to copy or rename ~/.mongorc.js to ~/.mongoshrc.js.
test> db.shutdownServer()
Browserslist: caniuse-lite is outdated. Please run:
npx browserslist@latest --update-db

Why you should do it regularly:
https://github.com/browserslist/browserslist#browsers-data-updating
MongoNetworkError: connection 1 to 127.0.0.1:27011 closed
test> exit

C:\Users\ZENBOOK>
```

Below this window is another smaller one titled 'C:\Windows\System32\cmd.exe' showing the command:

```
D:\mongodb\db\00_ServerTest>mongod -f node1.cfg
```

# Mongod Options

--help: output the various options for mongod with a description of their functionality.

- mongod --help or mongod -h

--dbpath <directory path>: Specify where all data files of the database are stored.

--port <port number>: specify the port on which mongod will listen for client connections.

- Run mongo shell connect to above mongod: mongosh --port 27018

--bind\_ip: specify which IP addresses mongod should bind to. When mongod binds to an IP address, clients from that address are able to connect to mongod.

- mongod --bind\_ip localhost --port 27018 --dbpath 'c:\mongoDB\data\db'
- mongod --bind\_ip localhost , 123.123.123.123 --port 27018 --dbpath 'c:\mongoDB\data\db'
- If using the bind\_ip option with external IP addresses, it's recommended to enable auth to ensure that remote clients connecting to mongod have the proper credential

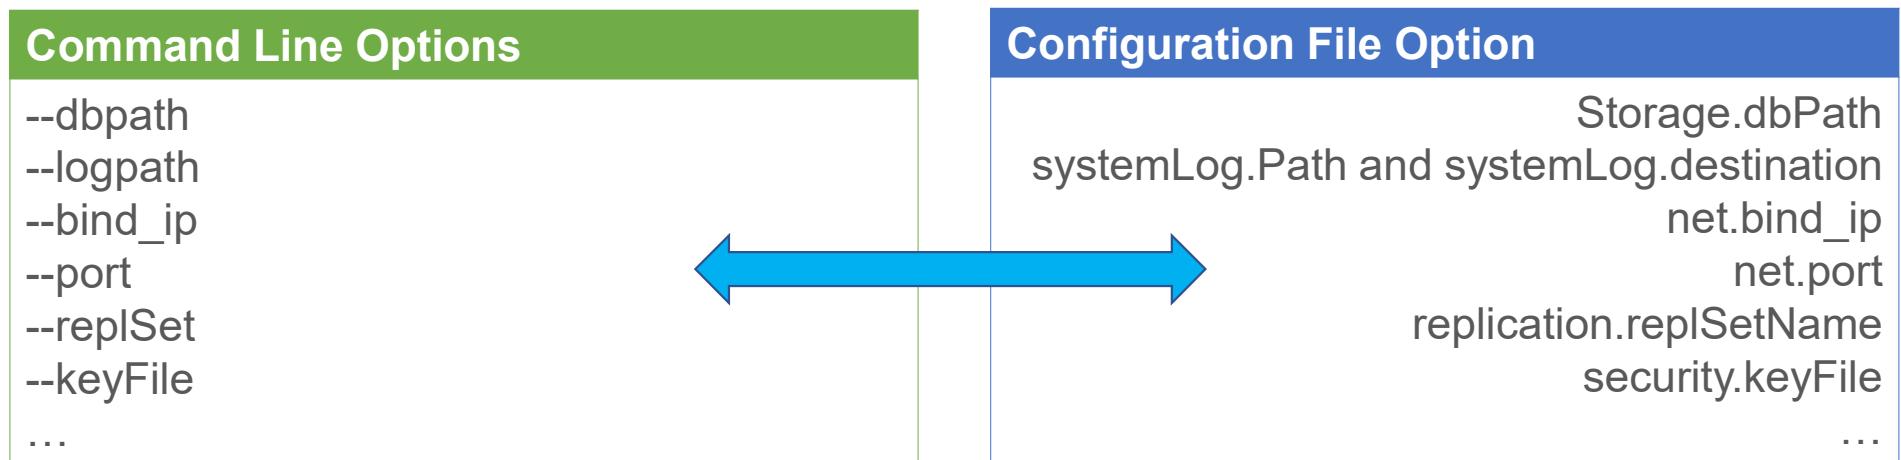
--auth: enables authentication to control which users can access the database. When auth is specified, all database clients who want to connect to mongod first need to authenticate.

# Mongod – Configuration File

- Configuration file is a way to organize the options you need to run the MongoD process into an easy to parse YAML (*Yet Another Markup Language*) file
- Why do we need to use configuration file?

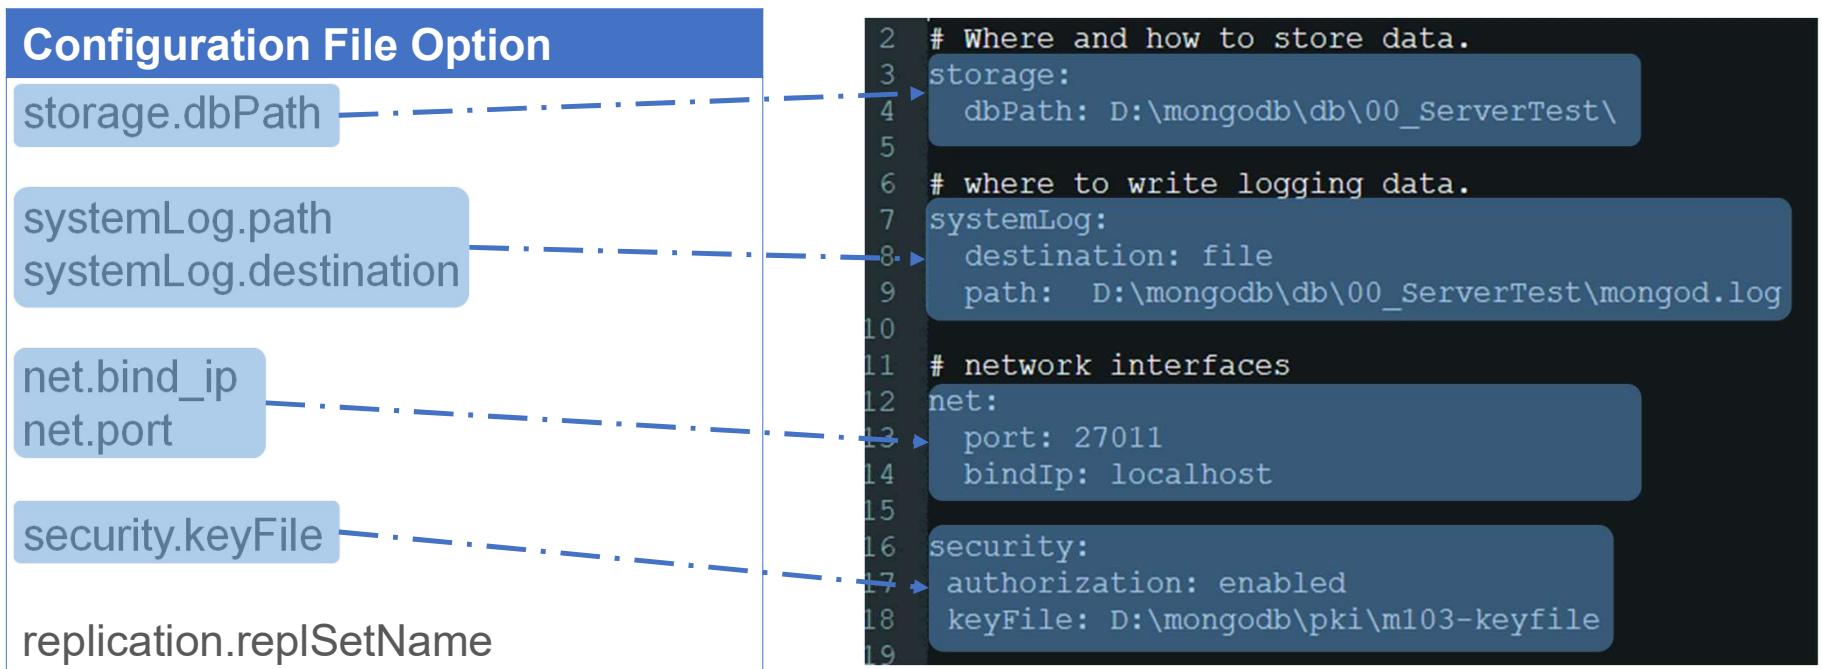
```
mongod --dbpath /data/db --logpath /data/log/mongod.log --replSet 'M103' --keyFile /data/keyfile --bind_ip '127.0.0.1'
```

# Mongod – Configuration File



# Mongod – Configuration File

YAML file



Launch mongod with the --config command line option:

```
mongod --config 'd:\CSDL_NoSQL\config_files\mongod.cfg'
```

# Mongod – Basic Commands

Cover a few of the basic commands necessary to interact with the MongoDB cluster.

These methods are available in the **mongodb shell** that wrap underlying database commands.

- **db.<method>()**: DB shell helpers, interact with the database.
- **db.<collection>.<method>()**: shell helpers for collection level operations.
- **rs.<method>()**: rs helper methods, control replica set deployment and management.
- **sh.<method>()**: sh helper mrthods, control sharded cluster deployment and management.

# Mongod – Basic Commands

User Management:

- `db.createUser()`
- `db.dropUser()`

Collection Management:

- `db.<collection>.renameCollection( <target>, <dropTarget> ) [dropTarget: optional]`
- `db.<collection>.createIndex( <keys>, <options>, <commitQuorum> )`
- `db.<collection>.drop( <options> )`

DB management:

- `db.dropDatabase( <writeConcern> ) [removes current database]`
- `db.createCollection( <name>, <options> )`

DB status:

- `db.serverStatus()`

# Mongod – Basic Security

Why do we have to secure the data?

## Authentication

- Verifies the identity of a user
- Answers the question : Who are you?

## Authorization

- Verifies the privileges of a user
- Answers the question: What do you have access to?

- **SCRAM**: default and most basic form of client authentication (password security)
  - **X.509**: certificate for authentication, more secure and more complex
  - **LDAP**
  - **Kerberos**
- Only for MongoDB Enterprise**

- Each user has one or more **Roles**.
- Each **Role** has one or more **Privileges**.
- A **Privilege** represent a group of **actions** and the **resources** that those actions apply to.

# Mongod – Basic Security

**Authorization:** Role Based Access Control

Roles support a high level of responsibility isolation for operational task:



- To enable role-based access control or authorization on cluster: enable authorization in configuration file (it implicitly enables authentication).
- By default, MongoDB doesn't give you any users.
- Always create a user with the administrative role first so you can create other users after.

```
#enable security  
security:  
    authorization: enabled
```

# Mongod – Basic Security

## Localhost Exception:

- Allows you to access a MongoDB server that enforces authentication **but** does not yet configured user for you to authenticate with.
- Must run mongo/mongosh from the **same host** running MongoDB server.
- Localhost exception **closes** after you create your first user.
- **Always** create a user with administrative privileges **first**.

# Mongod – Basic Security

## Example:

Run MongoDB server that enforces authentication (no user created)

```
mongod --config 'D:\HeQTCSDL_NoSQL\config_files\mongod.conf'
```

Run mongosh from the same host running MongoDB server

```
mongosh --host 127.0.0.1:27017
```

Create your first user

```
use admin
db.createUser( { user : 'root', pwd : 'root', roles : ['root'] })
```

Exit mongosh then run again with 'root' user

```
mongosh --username root --password root --authenticationDatabase admin
```

*or*

```
mongosh admin -u root -p root
```



```
net:
  port: 27017
  bindIp: 127.0.0.1
# enable security
security:
  authorization: enabled
```

# Mongod – Basic Security

## Roles in MongoDB

- **Build-In Roles:** Pre-packaged MongoDB Roles.
- **Custom Roles:** tailored roles to attend specific needs of specific users.
- **Database users:** will be granted roles to perform operations of MongoDB.

# Mongod – Basic Security

## Roles Structure

A role is composed of:

- Set of **privileges** that **role** enables
- All **privileges** that role defines will be made available to its users
- **Privilege** defines the **action**, or **actions**, that can be performed over a **resource**
- Resources:
  - Database
  - Collection or set of Collections
  - Cluster: Replica set, Shard Cluster

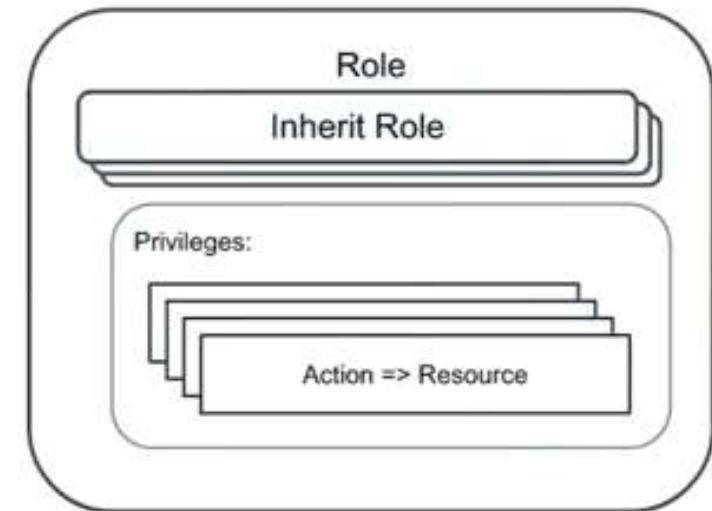
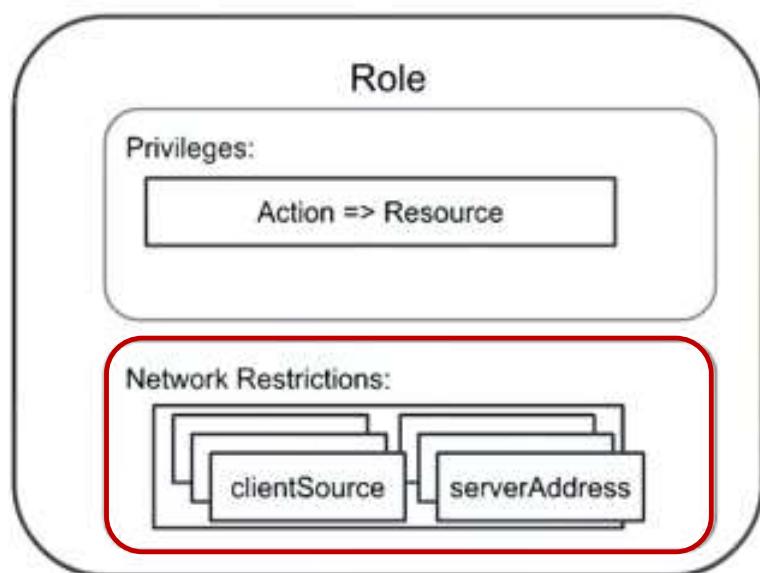
```
{resource: {cluster: true}, action: ['shutdown']}
```

A role with privilege, allowed to shut down any member of the cluster

# Mongod – Basic Security

## Roles Structure

A role can also inherit from other roles



We can also define network authentication restrictions at the role level

# Mongod – Basic Security

## Build-In Roles

| Role Levels             | Roles   |
|-------------------------|---|
| Database Users          | read, readWrite   |
| Database Administration | dbAdmin, userAdmin, dbowner   |
| Cluster Administration  | clusterAdmin, clusterManager, clusterMonitor, hostManager                         |
| Backup and Restore      | backup, restore   |
| Super User              | root (root is also a role at the all database level)                              |
| AllDatabase             | readAnyDatabase, readWriteAnyDatabase<br>dbAdminAnyDatabase, userAdminAnyDatabase |

*(read more Built-In Roles)*

# Mongod – Basic Security

## Build-In Roles: `userAdmin`

- Allows user to do all operations around user management. Not able to do anything related with data management or data modifications.
- Provides the ability to create and modify roles and users on the current database. Since the `userAdmin` role allows users to grant any privilege to any user, including themselves, the role also indirectly provides superuser access to either the database or, if scoped to the `admin` database, the cluster.

[\(read more `userAdmin` role\)](#)

## Example:

Run mongod with config file:

```
mongod --config 'D:\mongod.conf'
```

Run mongosh to connect to MongoDB server with root user:

```
mongosh admin -u root -p root
```

Create securityUser and grant `userAdmin` role

```
use admin          //all user should be created on the database admin for simplicity reasons
db.createUser( { user : 'securityUser', pwd : '123', roles : [ { db : 'admin', role : 'userAdmin' } ] } )
```

# Mongod – Basic Security

## Build-In Roles: dbAdmin

- Provides the ability to perform administrative tasks such as schema-related tasks, indexing, and gathering statistics. This role does not grant privileges for user and role management.
- Everything that is related with DDL (*data definition language*), this user will be able to perform.
- Everything that is related with the DML (*data modification language*) operations, he will not be able to do.

[\(read more dbAdmin role\)](#)

## Example:

Create securityUser and grant dbAdmin role

```
use admin
db.createUser( { user : 'DBAcourse', pwd : '123', roles : [ { db : 'mongoCourse', role : 'dbAdmin' } ] } )
//in this case, the roles of dbAdmin only be granted to mongoCourse db.
```

Roles can vary between databases. We can have a given user with different roles on a per database basis

```
db.grantRolesToUser( 'DBAcourse', [ { db : 'reporting', role : 'dbOwner' } ] )
```

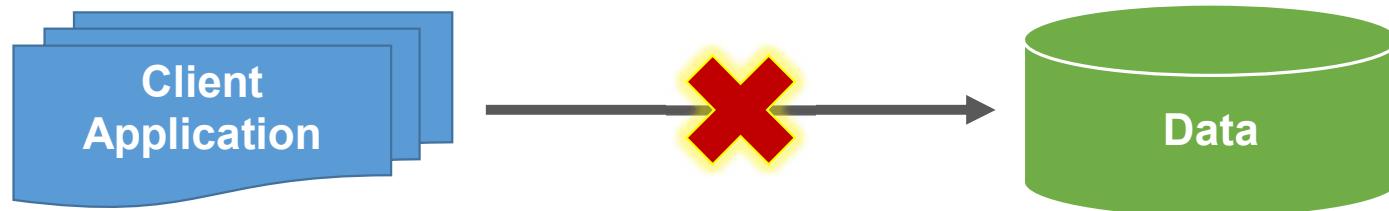
**dbOwner** role as a meta role. This role combines the privileges granted by the **readWrite**, **dbAdmin**, **userAdmin** roles

# 2. Replication

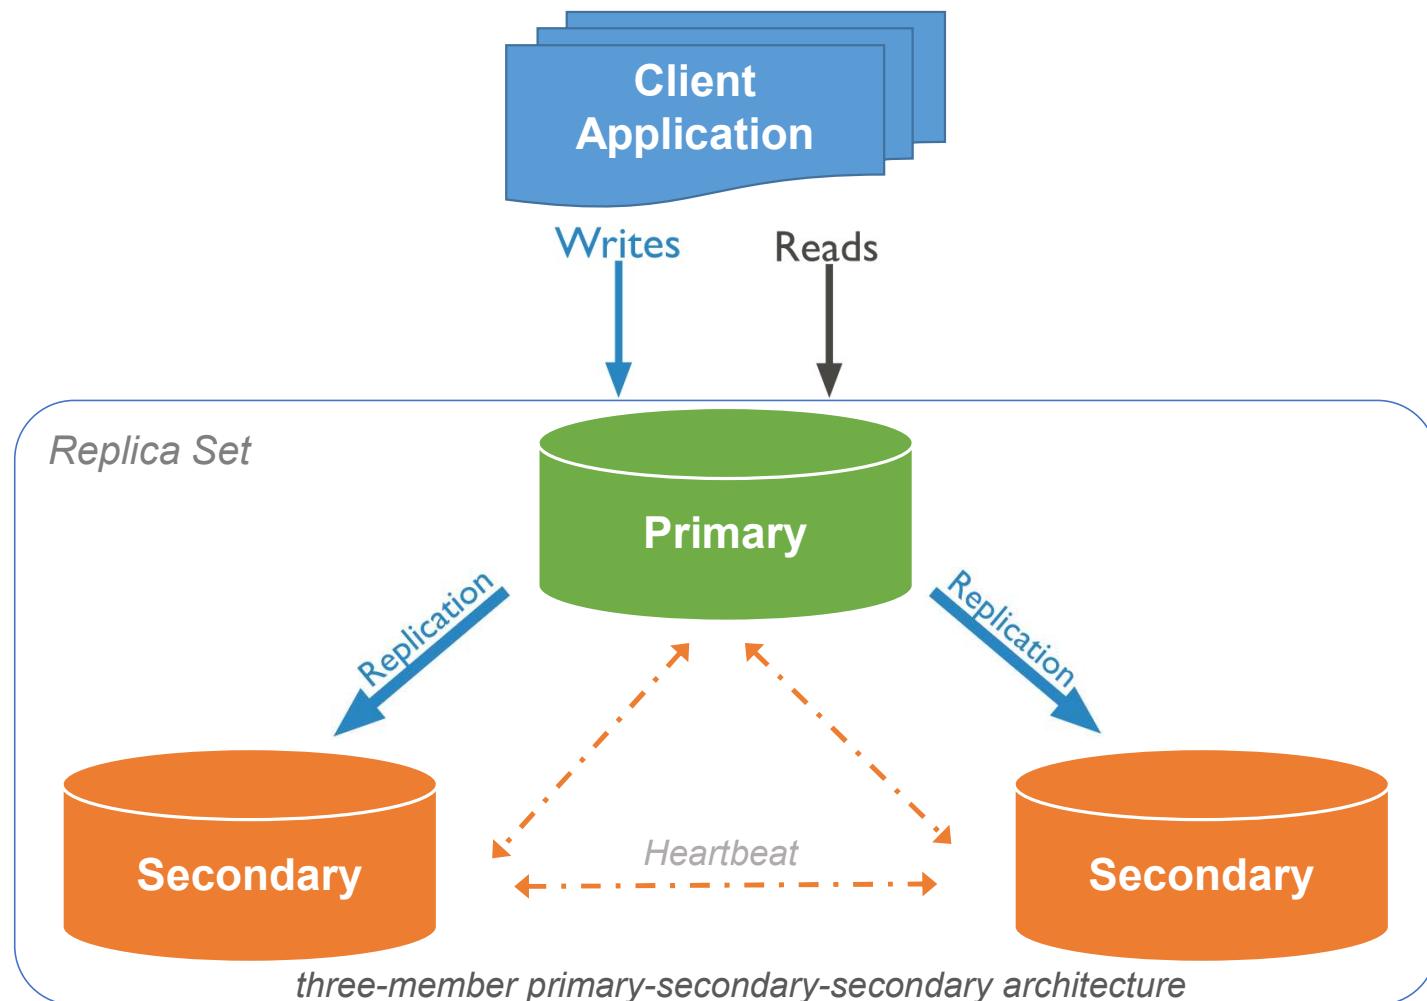
# Replication

- Replication: Maintain multiple copies of your data – **Really important**
- Why:
  - Can never assume all servers will always be available
  - To make sure, if server goes down, you can still access your data → **Redundancy and Data Availability**
  - Replication can provide increased read capacity as clients can send read operations to different servers

*If the database is hosted on a single server → **standalone node***

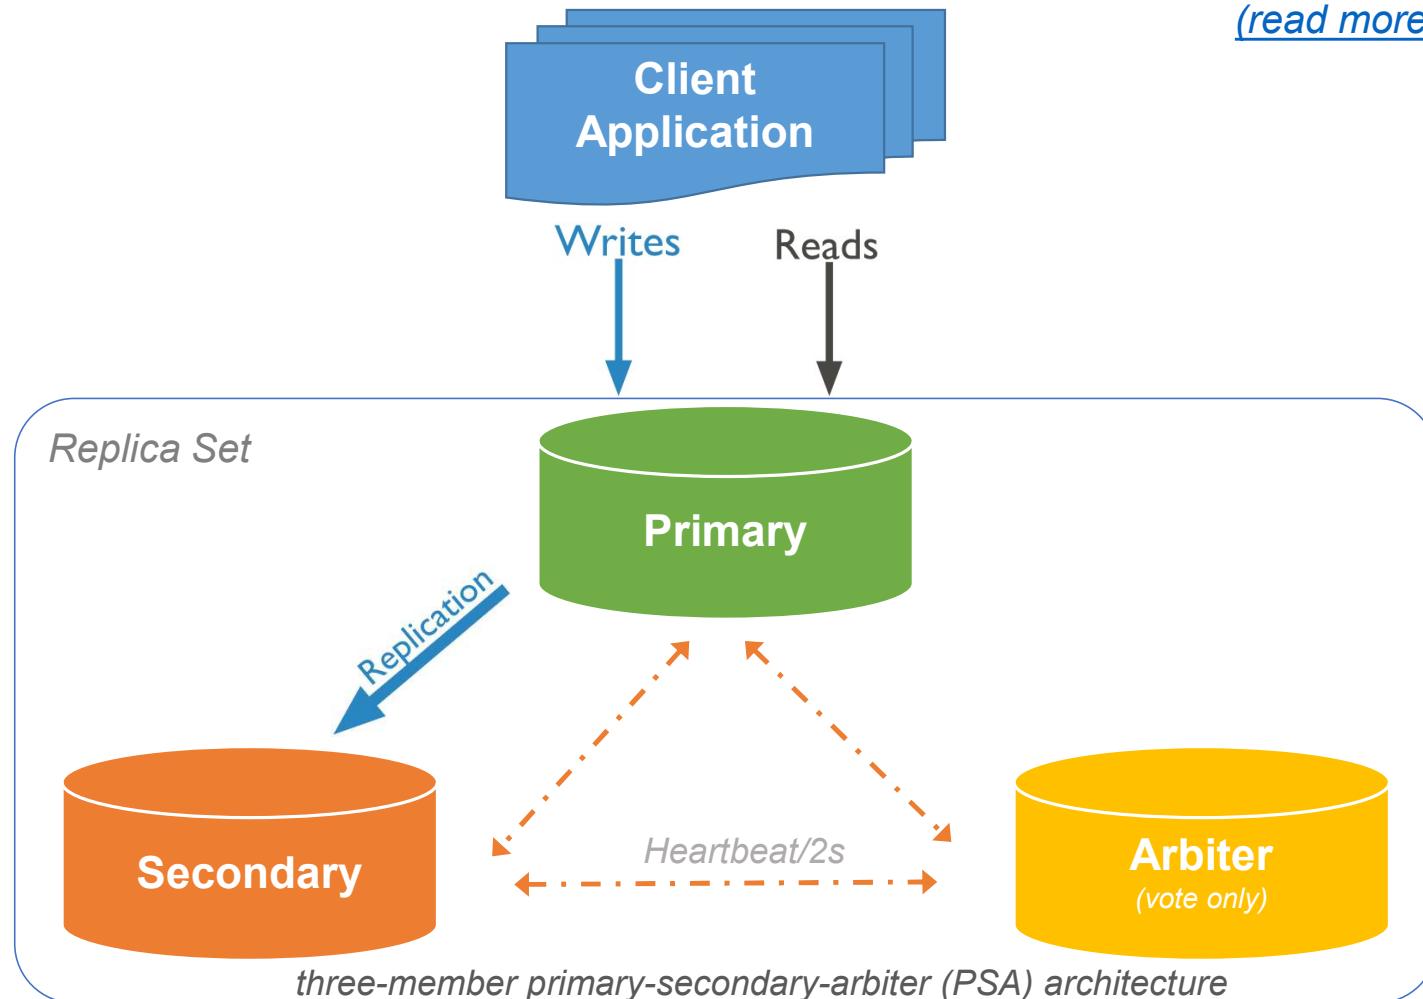


# Replication



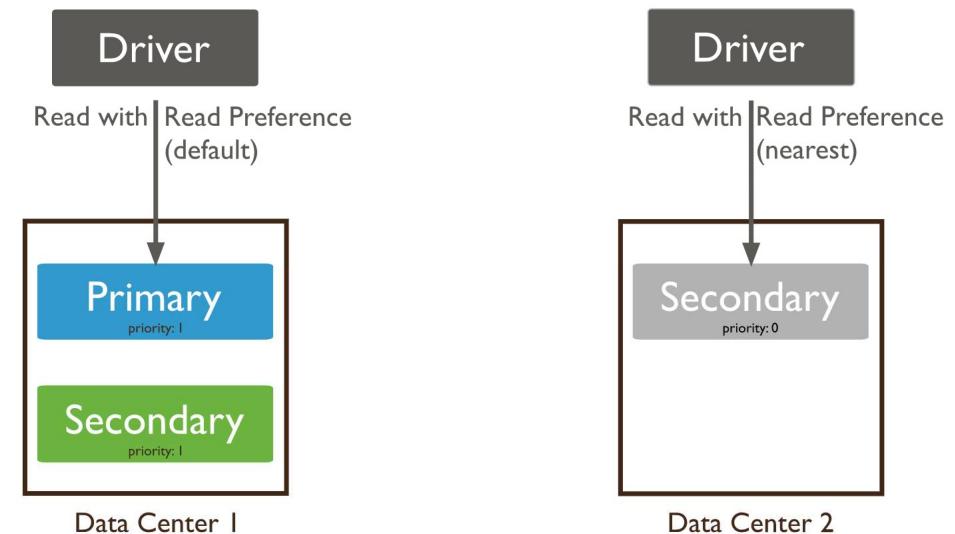
# Replication

[\(read more Replica Set Arbiter\)](#)



# Replication

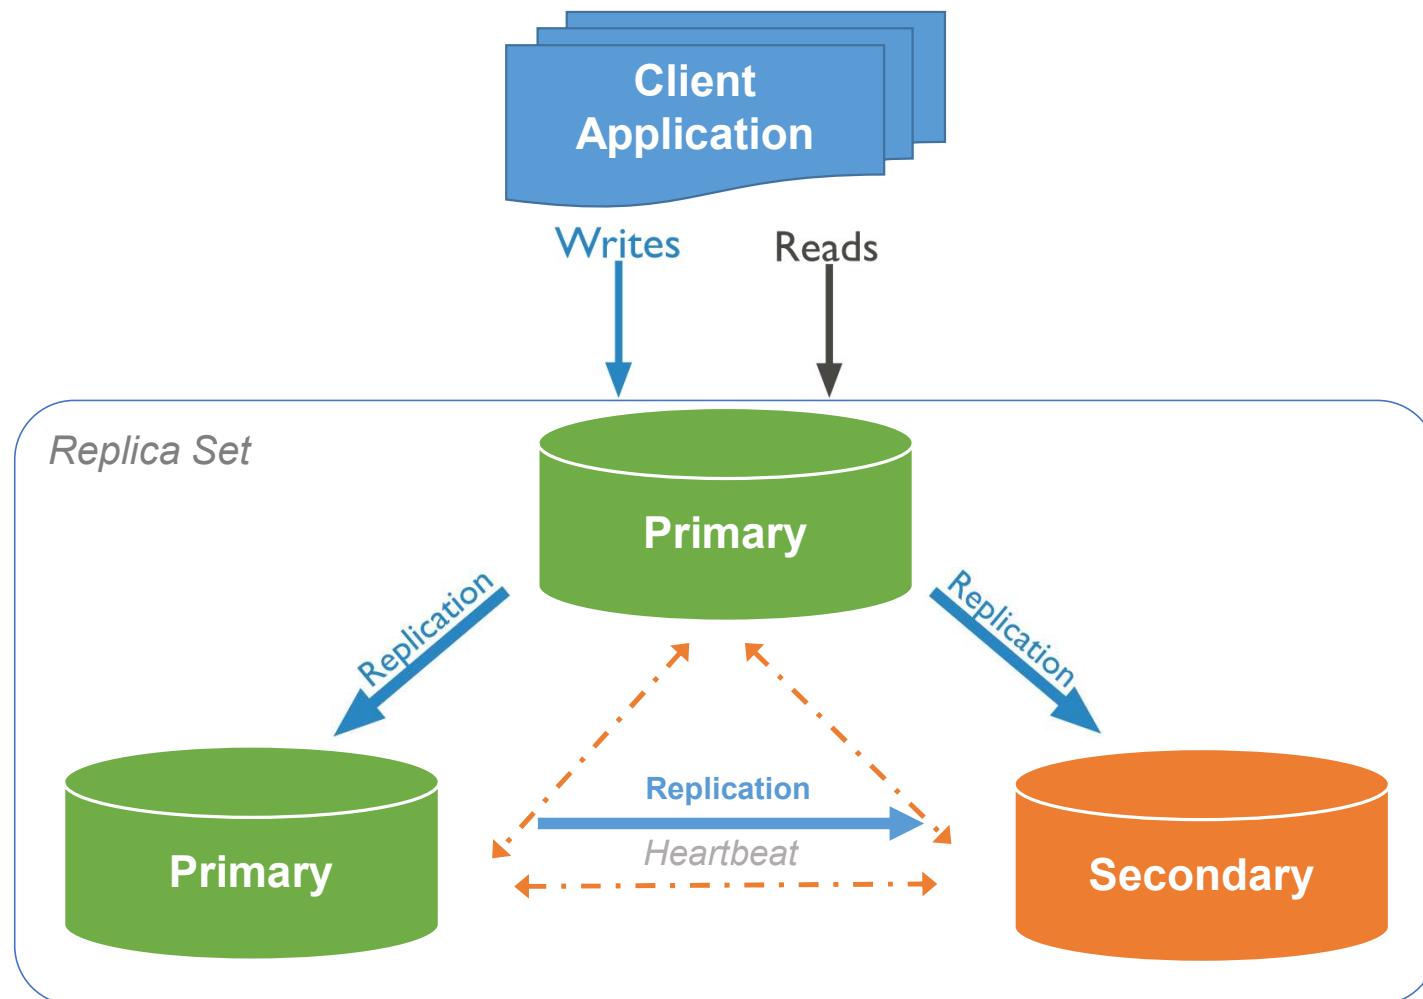
- A replica set is a group of mongod instances that maintain the same data set. A replica set contains several data bearing nodes and optionally one arbiter node. Of the data bearing nodes, one and only one member is deemed the primary node, while the other nodes are deemed secondary nodes.
- Although clients cannot write data to secondaries, clients can read data from secondary members. See [Read Preference](#) for more information on how clients direct read operations to replica sets.



- A secondary can become a primary. If the current primary becomes unavailable, the replica set holds an election to choose which of the secondaries becomes the new primary.

[\*\(read more Replica Set\)\*](#)

# Replication



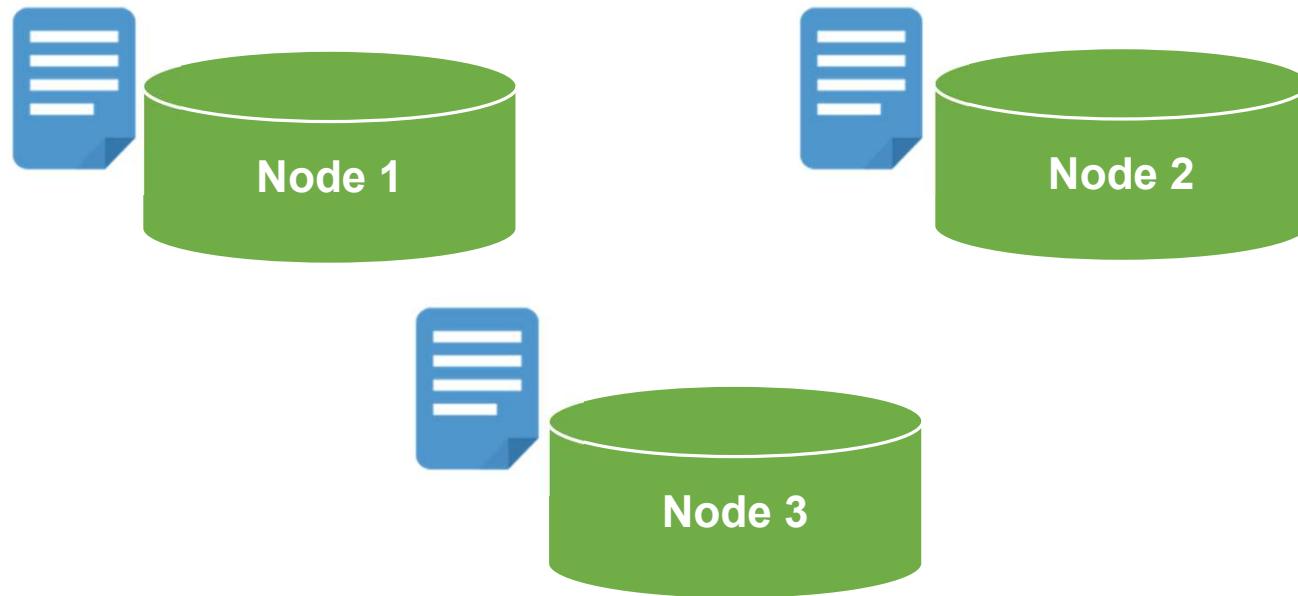
# Replication

- The replica set cannot process write operations until the election completes successfully. The replica set can continue to serve read queries if such queries are configured to run on secondaries.
- The median time before a cluster elects a new primary should not typically exceed 12 seconds. ([read more](#) [Replica Set Elections](#))
- You can configure a secondary member for a specific purpose. You can configure a secondary to:
  - Prevent it from becoming a primary in an election, which allows it to reside in a secondary data center or to serve as a cold standby. See [Priority 0 Replica Set Members](#).
  - Prevent applications from reading from it, which allows it to run applications that require separation from normal traffic. See [Hidden Replica Set Members](#).
  - Keep a running "historical" snapshot for use in recovery from certain errors, such as unintentionally deleted databases. See [Delayed Replica Set Members](#).

# Replication

## Setting up a Replica Set

- mongod won't be able to communicate with each other until we connect them



# Replication

## **Setting up a Replica Set:**

1. Use configuration file for standalone mongod;
2. Start a mongod with configuration file;
3. Start a mongo and connect to one of mongo instance;
4. Initialize replica set;
5. Create root user;
6. Exit out of this mongo and then log back in as m-admin user;
7. Add nodes to Replica set

# Replication

## Setting up a Replica Set:

1- Use configuration file for standalone mongod

|   |  |   |   |   |
|---|--|---|---|---|
| <b>storage:</b><br>dbPath: d:\mongodb\db\ <b>node1</b>  | <b>Optional, it is used to encrypt data exchanged between client application and mongodb</b> |  2 | d:\mongodb\db\ <b>node3</b>                                   |  |
| <b>net:</b><br>bindIp: localhost<br>port: <b>27011</b>  |  |   | localhost   |   |
| <b>security:</b><br>authorization: enabled<br>keyFile: d:\mongodb\pki\ <b>m-keyfile</b>                   | enabled<br><b>d:\mongodb\pki\m-keyfile</b>   |   | enabled<br><b>d:\mongodb\pki\m-keyfile</b>                    |   |
| <b>systemLog:</b><br>destination: file<br>path: d:\mongodb\db\ <b>node1\mongod.log</b><br>logAppend: true | file<br><b>d:\mongodb\db\</b> <b>node2\mongod.log</b><br>true                                |   | file<br><b>d:\mongodb\db\</b> <b>node3\mongod.log</b><br>true |   |
| <b>replication:</b><br>repSetName: <b>rep-example</b>   | <b>rep-example</b>   |   | <b>rep-example</b>  |   |

# Replication

## Setting up a Replica Set:

- 2- Start a mongod with configuration file:

```
mongod --config 'c:\mongoDB\configs\node1.conf'  
mongod --config 'c:\mongoDB\configs\node2.conf'  
mongod --config 'c:\mongoDB\configs\node3.conf'
```

- 3- Start a mongo and connect to one of mongo instance:

```
mongo --host 127.0.0.1:27011 || mongo --host localhost:27011
```

- 4- Initialize replica set:

```
rs.initiate()
```

- 5- Create root user:

```
use admin  
db.createUser({  
    user: 'm-admin',  
    pwd: 'm-pass',  
    roles: [ { role : 'root', db : 'admin' } ]  
})
```

# Replication

## Setting up a Replica Set

6- Exit out of this mongo and then log back in as m-admin user

```
mongo --host m-example/localhost:27011 -u m-admin -p m-pass --authenticationDatabase admin
```

7- Add nodes to Replica set

Replica set name

```
rs.add( 'localhost:27012' )  
rs.add( 'localhost:27013' )
```

To check status of Replica set: `rs.status()`

To check if the current node is primary: `rs.isMaster()`

```
rep-example [direct: primary] test> rs.isMaster()  
{  
    topologyVersion: {  
        processId: ObjectId("623355249e3d98501837545f"),  
        counter: Long("10")  
    },  
    hosts: [ 'localhost:27011', 'localhost:27012', 'localhost:27013' ],  
    setName: 'rep-example',  
    setVersion: 5,  
    ismaster: true,  
    secondary: false,  
    primary: 'localhost:27011',  
    me: 'localhost:27011',  
    electionId: ObjectId("7fffffff0000000000000001"),  
    lastWrite: {  
        opTime: { ts: Timestamp({ t: 1647532545, i: 1 }), t: Long("1") }  
    }  
}
```

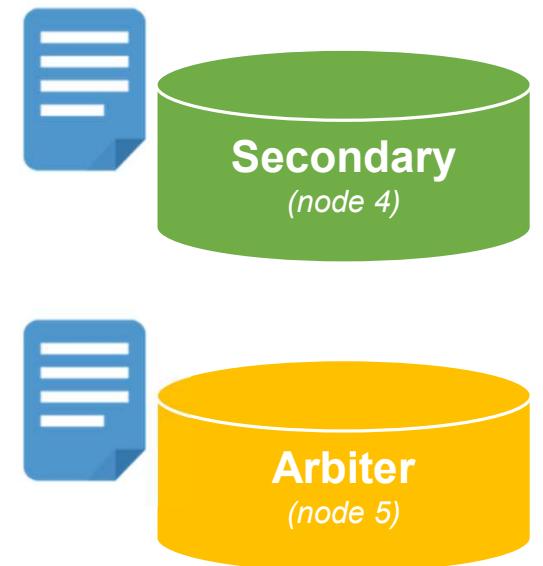
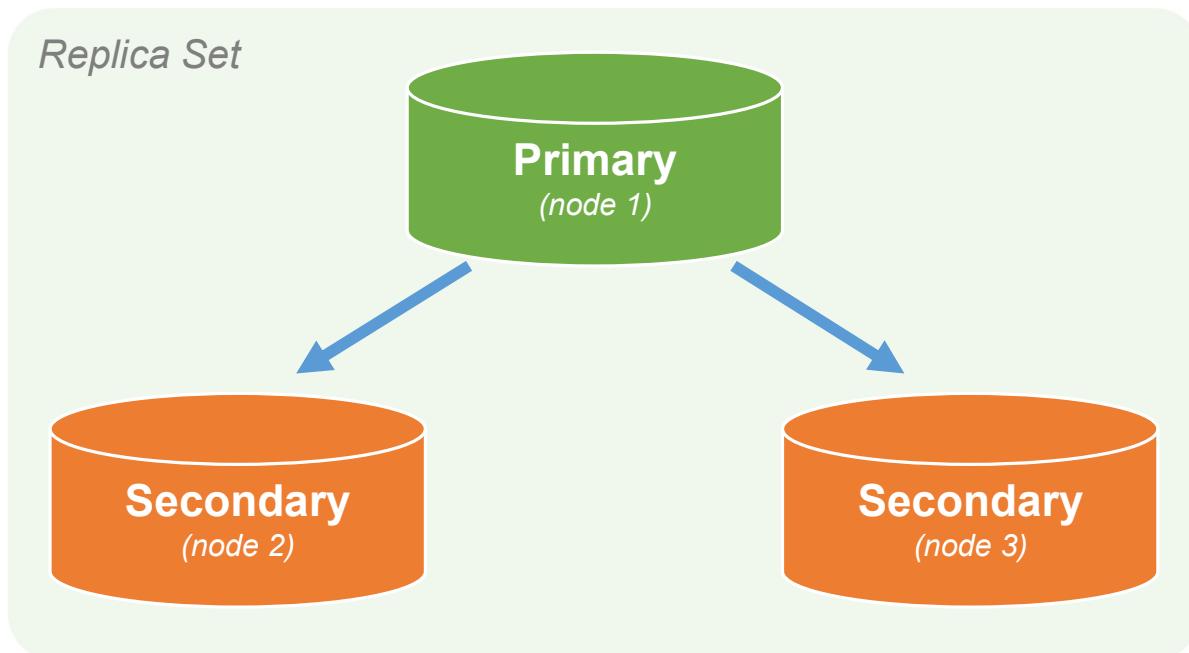
# Replication

## Replication Configuration Document:

- The replica set configuration document is a simple BSON document that we manage using a JSON representation
- Can be configured from the shell
- There are set of mongo shell replication helper methods that make it easier to manage
  - `rs.add()` : Adds a member to a replica set.
  - `rs.addArb()` : Adds an arbiter to a replica set.
  - `rs.initiate()` : Initializes a new replica set.
  - `rs.remove()` : Remove a member from a replica set.
  - `rs.reconfig()` : Reconfigures a replica set by applying a new replica set configuration object.
  - ... (*soft study*)

# Replication

Reconfiguring a Running Replica Set:



# Replication

## Reconfiguring a Running Replica Set:

1- Create config files for the secondaries 3 and arbiter nodes

```
storage:  
  dbPath: d:\mongodb\db\node4  
  
net:  
  bindIp: localhost  
  port: 27014  
  
security:  
  authorization: enabled  
  keyFile: d:\mongodb\pki\m-keyfile  
  
systemLog:  
  destination: file  
  path: d:\mongodb\db\node4\mongod.log  
  logAppend: true  
  
replication:  
  replSetName: rep-example
```



```
storage:  
  dbPath: d:\mongodb\db\arbiter  
  
net:  
  bindIp: localhost  
  port: 28000  
  
security:  
  authorization: enabled  
  keyFile: d:\mongodb\pki\m-keyfile  
  
systemLog:  
  destination: file  
  path: d:\mongodb\db\arbiter\mongod.log  
  logAppend: true  
  
replication:  
  replSetName: rep-example
```



# Replication

## Reconfiguring a Running Replica Set:

- 2- Starting up mongod processes for our fourth node and arbiter

```
mongod --config 'c:\mongoDB\configs\node4.conf'
```

```
mongod --config 'c:\mongoDB\configs\arbiter.conf'
```

- 3- Run Mongo shell and connect to the replica set m-example

```
mongo --host m-example/localhost:27011 -u 'm-admin' -p 'm-pass' --authenticationDatabase 'admin'
```

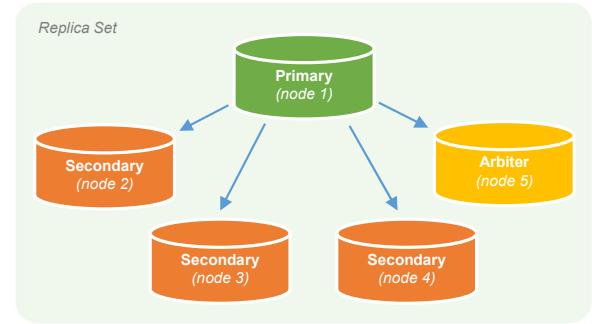
- 4- From the Mongo shell of the replica set, adding the new secondary and the new arbiter:

```
rs.add('localhost:27014')
```

```
rs.addArb('localhost:28000')
```

- 5- Checking replica set make up after adding two new nodes:

```
rs.isMaster()
```



```
rep-example [direct: primary] test> rs.isMaster()
{
  topologyVersion: {
    processId: ObjectId("6233fee5d005505fb75376ed"),
    counter: Long("9")
  },
  hosts: [
    'localhost:27011',
    'localhost:27012',
    'localhost:27013',
    'localhost:27014'
  ],
  arbiters: [ 'localhost:28000' ],
  setName: 'rep-example',
  setVersion: 8,
  ismaster: true,
  secondary: false,
  primary: 'localhost:27011',
  me: 'localhost:27011',
  selectionId: ObjectId("7FFFFFFFFFFF0000000000000002")
}
```

# Replication

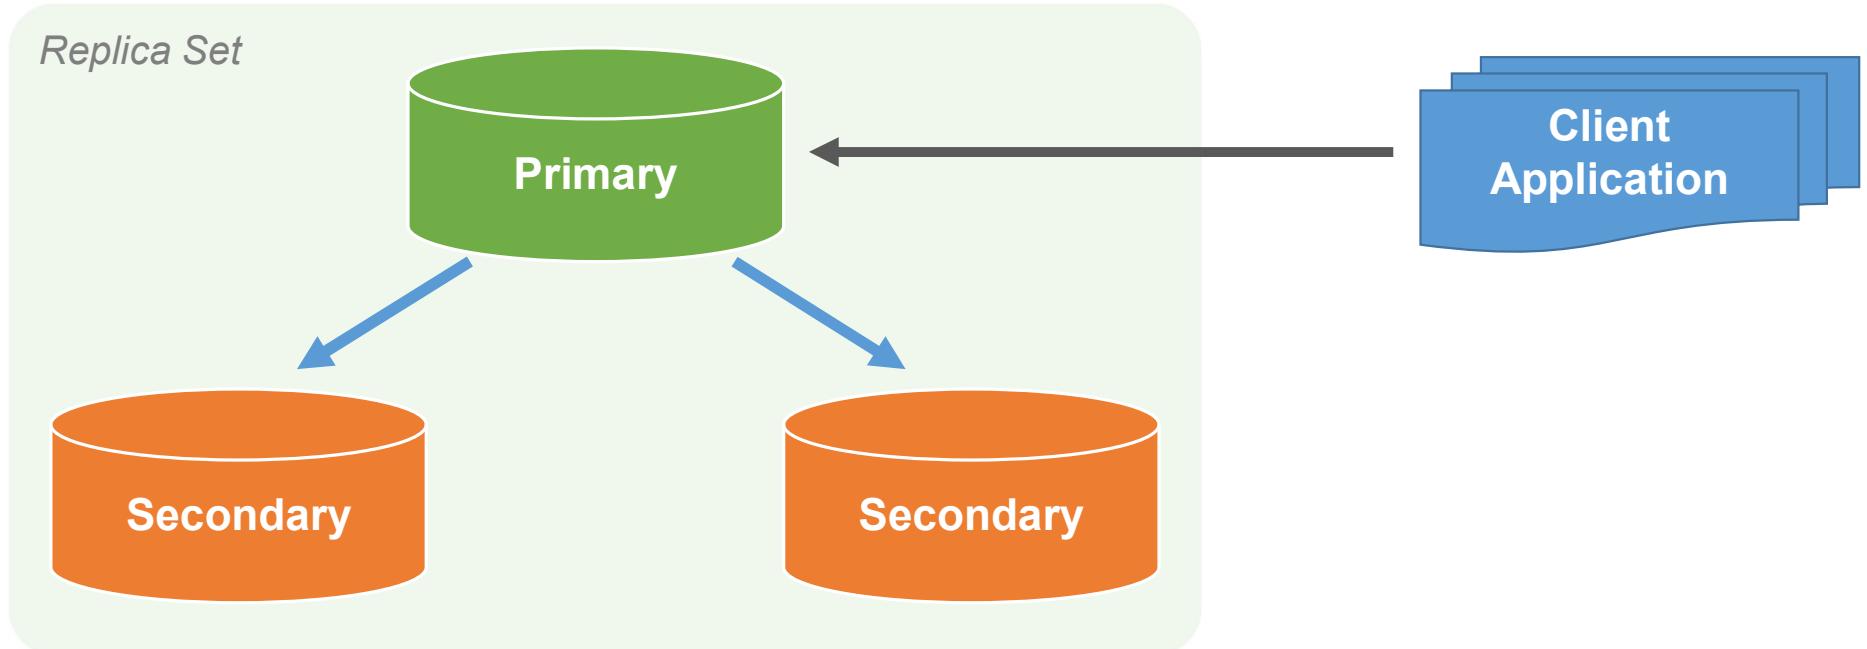
## Reconfiguring a Running Replica Set:

- Removing the arbiter from our replica set:  
`rs.remove( 'localhost:28000' )`
- Assigning the current configuration to a shell variable we can edit, in order to reconfigure the replica set:  
`cfg = rs.conf()`
- Editing our new variable cfg to change topology - specifically, by modifying cfg.members:  
`cfg.members[3].votes = 0`  
`cfg.members[3].hidden = truea`  
`cfg.members[3].priority = 0`
- Updating our replica set to use the new configuration cfg:  
`rs.reconfig( cfg )`

# Replication

## Failover and Elections:

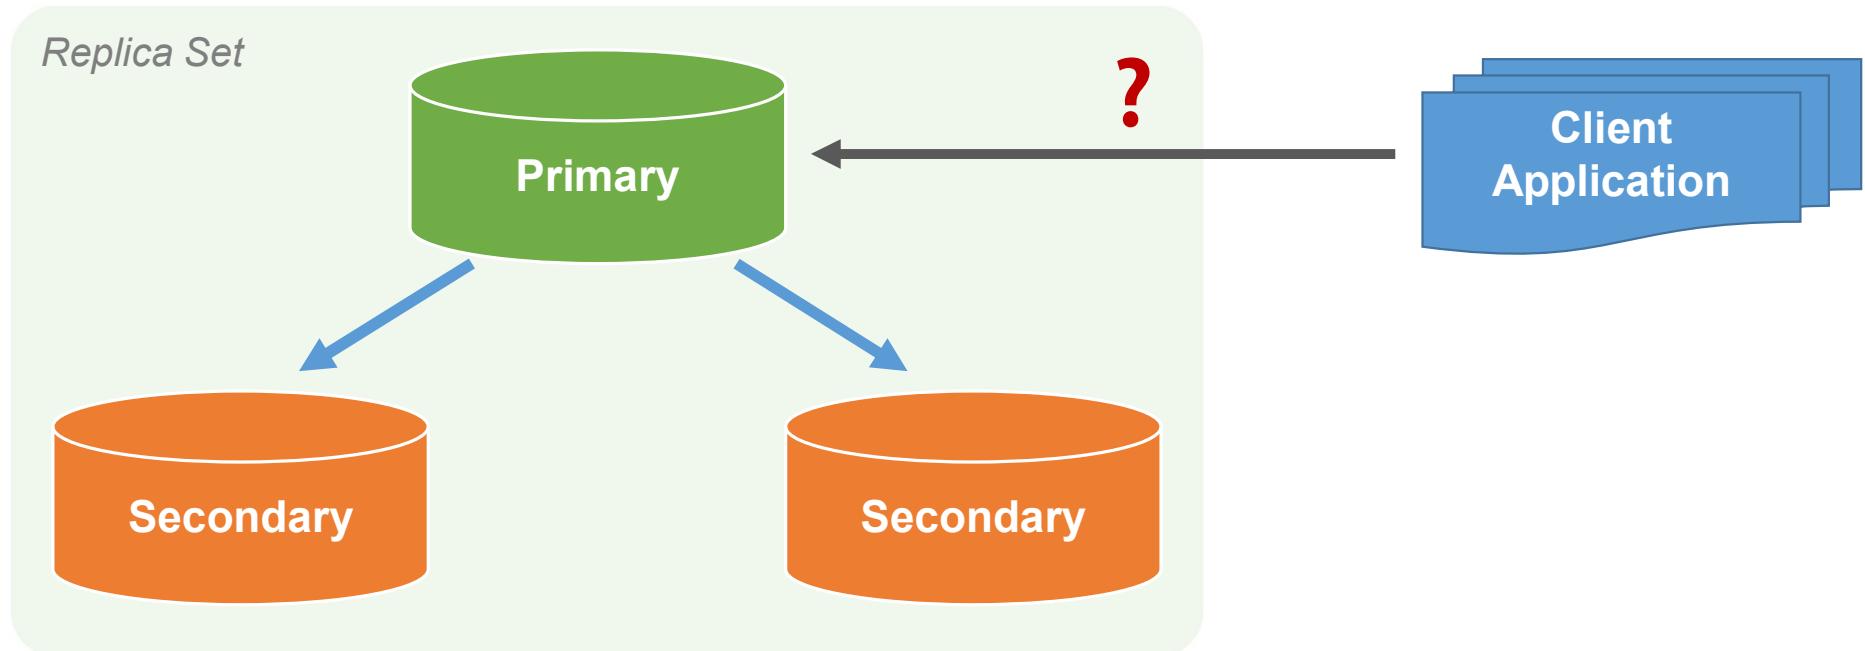
- Primary node is the first point where the client application accesses the database.
- if secondaries go down, the client will continue communicating with the node acting as primary until the primary is unavailable.



# Replication

## Failover and Elections:

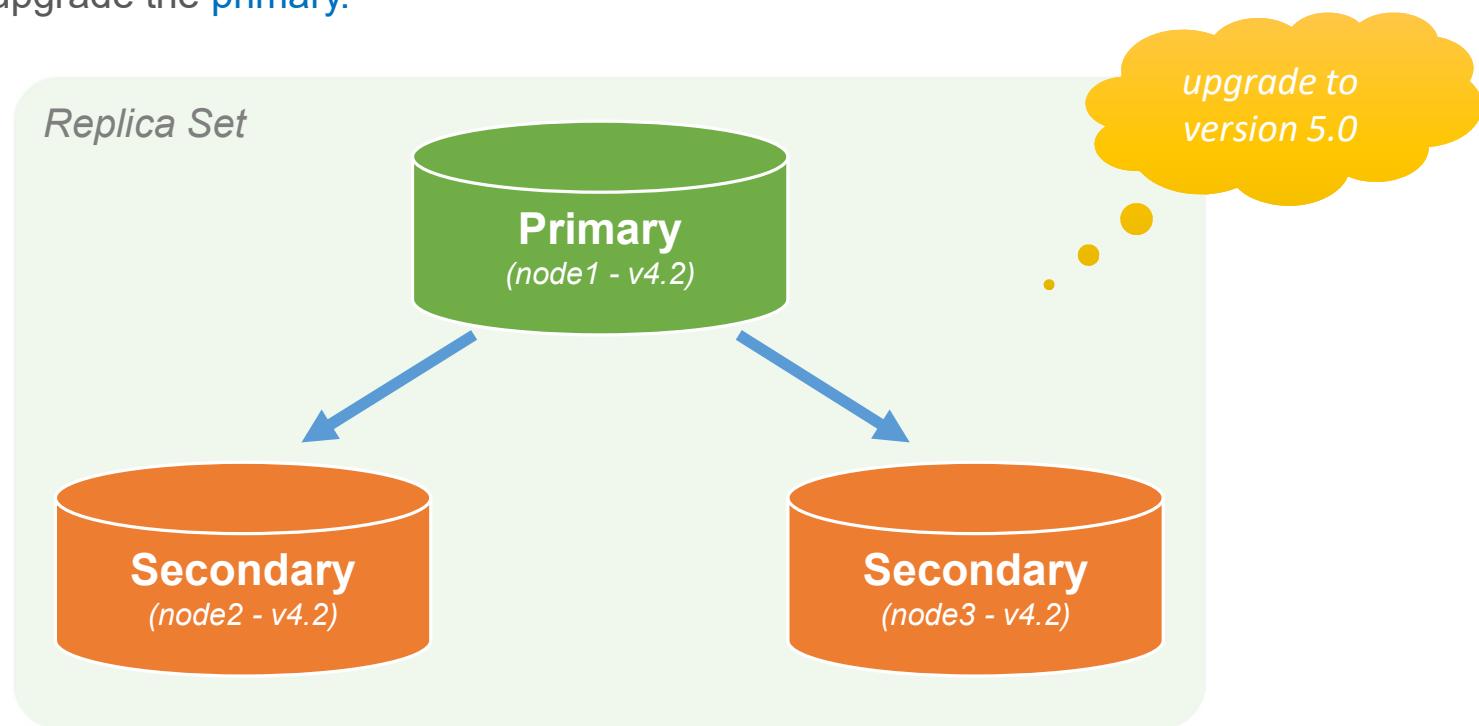
- What would cause a primary to become unavailable? → a common reason is maintenance.



# Replication

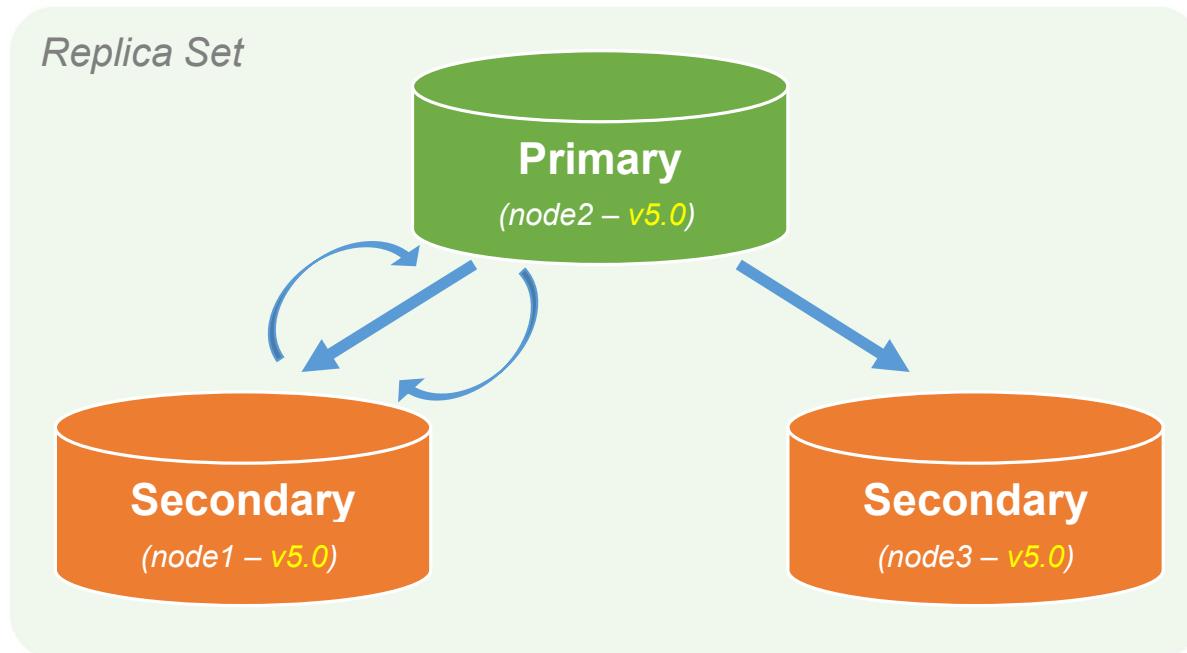
## Failover and Elections:

- Let's say we want to **roll upgrade** on a three nodes replica set.
- A rolling upgrade just means we're upgrading **one server at a time**, starting with the **secondaries** and eventually, we'll upgrade the **primary**.



# Replication

Failover and Elections:



# 3. Sharding

# Mongod – Sharding

## What is Sharding?

- In a replica set, we have more than one server in our database and each server has to contain the entire dataset
- What do we do when the data grows, and the servers can't work properly?

**Vertical Scaling:** increase the capacity of individual server:

- More RAM
- More disk space
- More powerful CPU



- Potentially become very expensive
- Cloud-based providers aren't going to let us scale vertically forever

